

Enterprise System Health Monitoring & ML Prediction Framework

Executive Overview

Transform your current laptop monitoring tool into a **distributed network monitoring system** with predictive analytics, all operating completely offline within DRDO's secure infrastructure.

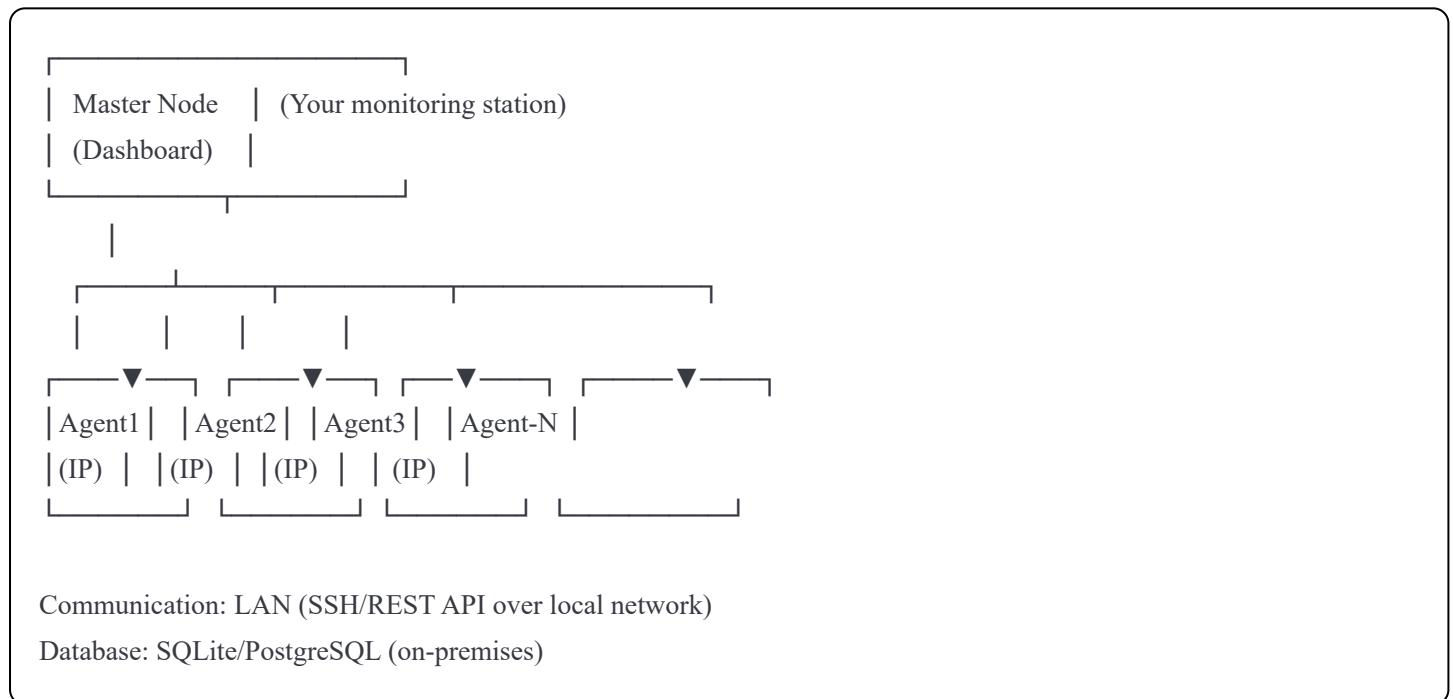
PHASE 1: Multi-Device Monitoring Architecture

1.1 Network Discovery & Remote Monitoring

Current State: Monitoring single system

Target State: Monitor 50-1000+ systems on DRDO network

Implementation Strategy:



Key Components:

- **Master Dashboard:** Central aggregation point running on secured workstation
- **Agent Software:** Lightweight Python service deployed on each device (via group policy/automated deployment)
- **Communication Protocol:** SSH tunneling or simple REST API over HTTPS (with self-signed certs)
- **Offline Sync:** Store metrics locally, sync when available

1.2 Remote Data Collection Methods

Option A: SSH-Based Collection (Most Secure for DRDO)

```

python

# Pseudo-code structure

class RemoteSystemMonitor:
    def __init__(self, ip, username, ssh_key_path):
        self.ssh_client = paramiko.SSHClient()
        self.ip = ip

    def get_remote_metrics(self):
        # Execute psutil commands remotely
        # Return JSON with metrics
        pass

```

Advantages:

- No additional ports needed
- Encryption built-in
- Credential management via SSH keys
- DRDO network compatible

Deployment:

- Deploy agent Python script on each device
- Set up cron jobs (Linux) or Task Scheduler (Windows)
- Agent runs every 5 minutes, logs metrics locally
- Master pulls data via SSH every 15 minutes

Option B: Lightweight Agent with Local Storage

Each Device:

1. Runs lightweight Python agent (background service)
2. Collects metrics every 5 minutes
3. Stores in local SQLite DB
4. Master pulls via API/SSH on schedule

1.3 Device Discovery

Automated Network Scanning:

1. Get network subnet from DRDO network config
2. Use nmap/ARP scanning to find active devices
3. Maintain inventory database
4. Alert on new devices/offline devices

PHASE 2: Time-Series Data Collection & Storage

2.1 Historical Data Model

SYSTEMS TABLE:

- system_id (unique)
- ip_address
- hostname
- os_type
- location (building/lab)
- department
- added_date

METRICS TABLE (Time-Series):

- timestamp
- system_id
- cpu_avg, cpu_max, cpu_cores_usage[]
- memory_used, memory_total, swap
- disk_usage (by partition)
- disk_io (read/write speeds)
- network_io (sent/received)
- temperature
- battery_status
- process_count
- network_connections

2.2 Data Retention Strategy

- Real-time: Keep 7 days (5-min intervals)
Daily: Keep 6 months (1 daily aggregate)
Monthly: Keep 2 years (1 monthly aggregate)
Archival: Compressed backups for compliance

Offline Storage:

- Use **SQLite** for local node storage
- Use **PostgreSQL** for master database (optional, can run offline)
- CSV exports for archival

- JSON for lightweight transfer
-

PHASE 3: Machine Learning - Predictive Analytics

3.1 Problem Statements ML Can Solve

1. Hardware Failure Prediction (6-month outlook)

Input Features (30-day window):

- CPU avg/max/variance
- Memory peak usage trends
- Disk I/O patterns
- Temperature trend
- System uptime/reboot frequency
- Hard disk SMART data (if available)

Output: Probability of failure in next 6 months

Models: Random Forest, Gradient Boosting, LSTM

Example Prediction:

- Device X: 87% chance of disk failure (increasing bad sectors detected)
- Device Y: 12% chance of failure (all metrics stable)

2. Performance Degradation Detection

Detect: Is performance declining over time?

- CPU bottlenecks emerging
- Memory leaks in services
- Disk becoming fragmented
- Network issues arising

Output: Anomaly score + recommendation

3. Capacity Planning

Forecast: When will capacity reach critical?

- Memory trending 5% up per month → critical in 8 months
- Disk usage pattern suggests overflow in 3 months
- Network saturation prediction

4. Abnormal Behavior Detection

Real-time anomaly detection:

- CPU spike from 10% to 95% (malware/runaway process?)
- Memory leak (gradual increase, never released)
- Unexpected network traffic
- Unusual process count

5. Energy Consumption Prediction

For laptop devices:

- Battery health degradation
- Usage pattern forecasting
- Charging time predictions

3.2 ML Approach for DRDO (Offline)

Architecture:

DATA COLLECTION (15 days minimum)



FEATURE ENGINEERING



OFFLINE MODEL TRAINING (Weekly)



MODEL STORAGE (SQLite/joblib)



REAL-TIME INFERENCE (Every 30 mins)



ALERTS & RECOMMENDATIONS

Model Types:

A. Time-Series LSTM (Best for trends)

Predicts: Next 30 days CPU/Memory/Disk usage

Gives: Early warning if trending toward critical

B. Isolation Forest (Anomaly Detection)

Detects: Unusual metric combinations

Real-time with zero network dependency

Lightweight, runs on each device

C. Gradient Boosting (Classification)

Predicts: Failure probability (binary: fail/no-fail)

Training: Historical data from all devices

Labels: Devices that actually failed (root cause analysis post-incident)

D. Unsupervised Clustering

Groups: Similar devices together

Identifies: Which device cohorts fail together

Action: Preventive maintenance on similar systems

3.3 Training Data Strategy

Initial Bootstrap (Without Historical Data):

Week 1-2: Collect baseline metrics (all devices)

Week 3-4: Collect during normal operations

Week 5-6: Collect during stress tests

Create synthetic scenarios:

- Simulate disk filling up
- CPU stress tests
- Memory allocation tests
- Network saturation
- Hardware degradation patterns

Label data:

- Known failures in DRDO network history
- Current device age/performance curve
- Manufacturer reliability curves

Continuous Learning:

Every incident:

1. Root cause analysis
2. Extract features leading to failure
3. Add as training example
4. Retrain model (weekly)
5. Improve predictions

Example: If 5 devices failed with:

- Temp > 85°C + CPU usage > 80% + Uptime < 30 days
- Model learns this pattern predicts failure

3.4 Feature Engineering for Better Predictions

Raw Metrics → Features:

1. STATISTICAL FEATURES:

- 7-day rolling average (CPU, Memory)
- Standard deviation (volatility)
- Max/Min/Range
- Rate of change (derivative)

2. TREND FEATURES:

- Linear regression slope (is it increasing?)
- Exponential growth rate
- Cyclical patterns (office hours vs night)

3. DOMAIN FEATURES:

- Thermal stress index: (Current Temp / Max Temp) * Duration
- Memory leak indicator: Memory(t) - Memory(t-24h)
- Disk fragmentation risk: (Used Space / Total Space) ^ 2
- Device age in days
- OS version maturity
- Running process count

4. INTERACTION FEATURES:

- CPU * Temperature (thermal load stress)
- Memory * Swap (pressure indication)
- Disk Usage * IO Speed (performance degradation)

PHASE 4: Implementation Roadmap

Sprint 1 (Week 1-2): Multi-Device Architecture

- Build SSH client in Python (paramiko)
- Implement agent deployment script
- Create centralized metrics collector
- Set up PostgreSQL/SQLite infrastructure

Sprint 2 (Week 3-4): Dashboard Expansion

- Add device inventory management
- Create network topology visualization
- Build time-series graphing (Plotly/Recharts)
- Implement alert threshold system

Sprint 3 (Week 5-6): Data Pipeline

- ETL pipeline for metric aggregation
- Data cleaning & validation
- Implement time-series compression
- Create backup/archival system

Sprint 4 (Week 7-8): ML Foundation

- Historical data collection (minimum 4 weeks)
- Feature engineering pipeline
- Train baseline models (Random Forest, Isolation Forest)
- Create model evaluation framework

Sprint 5 (Week 9-10): Predictive Features

- Deploy failure prediction model
- Build anomaly detection
- Create trend forecasting
- Develop alert generation

Sprint 6 (Week 11-12): Optimization & Hardening

- Model optimization (quantization for edge)
 - Performance tuning
 - Security hardening (encryption, audit logs)
 - Documentation & training
-

PHASE 5: Advanced Features for Enterprise

5.1 Security-Focused Monitoring

1. PROCESS MONITORING:

- Track unauthorized processes
- Detect suspicious activity patterns
- Monitor network connections by process
- Alert on unusual service startups

2. PATCH MANAGEMENT:

- Track OS/software versions
- Identify missing critical patches
- Correlate with known vulnerabilities
- Recommend update schedules

3. COMPLIANCE TRACKING:

- Log all monitoring access

- Audit trail of configuration changes
- User login tracking
- Incident investigation support

4. ENCRYPTION:

- All inter-device communication: TLS 1.3
- Database encryption: AES-256
- Backup encryption
- Data sanitization policies

5.2 Compliance & Reporting

Automated Reports:

1. DAILY ALERTS:

- Devices offline
- Critical thresholds exceeded
- Anomalies detected
- Predictive failures (next 7 days)

2. WEEKLY SUMMARY:

- Device utilization trends
- Incident summary
- Maintenance recommendations
- Network health score

3. MONTHLY EXECUTIVE REPORT:

- Infrastructure health (0-100 score)
- Failure predictions
- Capacity planning recommendations
- ROI on preventive maintenance
- Cost projections

4. COMPLIANCE REPORT:

- Security audit trail
- Access logs
- Configuration changes
- Patch compliance status

5.3 Predictive Maintenance Scheduling

System: "Device X will fail in 14 days"

Action: Automatically generate maintenance ticket

Maintenance Scheduler:

1. Predict failure window

2. Check device criticality
3. Check device availability schedule
4. Find optimal maintenance window
5. Generate ticket for technician
6. Track maintenance completion
7. Verify device performance post-maintenance

PHASE 6: Complete Offline Architecture

6.1 Offline Requirements

- ✓ Zero internet dependency
- ✓ All data stays on DRDO network
- ✓ Models train on-premises
- ✓ Predictions run locally
- ✓ Complete audit trail
- ✓ Air-gap safe

6.2 Technology Stack (Fully Offline)

FRONTEND:

- PyQt5 / Tkinter (offline GUI)
- Web dashboard with offline caching

BACKEND:

- FastAPI (local REST API)
- PostgreSQL / SQLite (local database)

DATA SCIENCE:

- scikit-learn (models)
- XGBoost / LightGBM (gradient boosting)
- TensorFlow Lite (if using LSTM)

DEPLOYMENT:

- Docker (containerize entire stack)
- Local registry (images stored locally)

MONITORING:

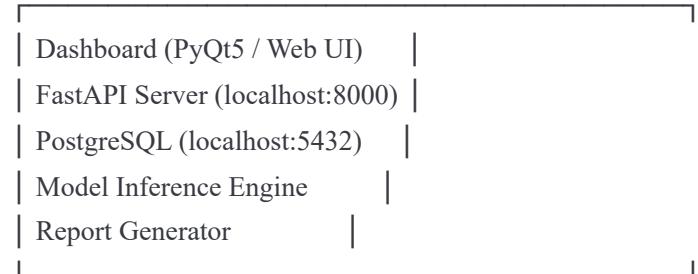
- Prometheus + Grafana (optional, or use Plotly)
- Custom dashboards

COMMUNICATION:

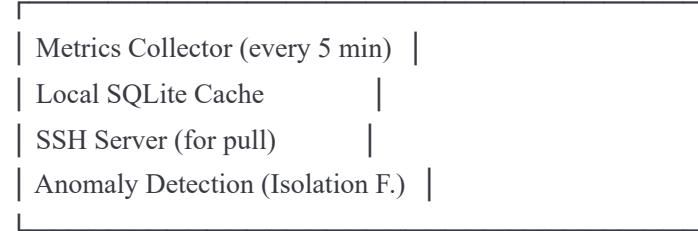
- SSH (over LAN only)
- HTTPS with self-signed certificates

6.3 Deployment Architecture

MASTER NODE (Central Server):



AGENT NODES (Each Device):



NETWORK (DRDO LAN only):

SSH Tunnels + Self-signed TLS

Port 22 (SSH) + Port 8000 (API)

No external connectivity required

PHASE 7: Data Examples & Expected Outputs

7.1 Failure Prediction Example

INPUT (30-day metrics for Device-C):

- CPU Usage: avg 45%, max 98%, trend +3% per week
- Memory: avg 60%, max 85%, trend stable
- Disk: 78% full, growing 2% per week
- Temp: avg 62°C, max 79°C, trend +1.5°C per week
- Uptime: 89 days (stable)
- Age: 3.2 years

MODEL OUTPUT:

Failure Risk Score: 73%

Primary Risk: Disk capacity (will be full in 5 weeks)

Secondary Risk: Thermal degradation

Tertiary Risk: Hardware age (3.2 years, typical lifespan 5 years)

RECOMMENDATION:

- Priority: HIGH
- Action: Archive data to external storage within 2 weeks

- Backup: Increase backup frequency to daily
- Timeline: Plan replacement for month 6
- Cost: Server capacity upgrade needed

7.2 Anomaly Detection Example

ALERT: Device-F unusual behavior detected

Metric	Normal Range	Current	Anomaly Score
CPU Usage	15-35%	87%	CRITICAL
Memory Usage	40-60%	76%	HIGH
Network Outbound	0-5 MB/s	45 MB/s	CRITICAL
Process Count	120-150	512	CRITICAL
Temperature	45-60°C	73°C	MEDIUM

Likely Causes:

1. Malware/Ransomware activity (45% confidence)
2. Misconfigured backup process (35% confidence)
3. Resource leak (15% confidence)
4. Stress test in progress (5% confidence)

Recommended Actions:

1. IMMEDIATE: Check process manager for suspicious processes
2. VERIFY: Check with team if stress test is scheduled
3. SCAN: Run antivirus scan if cleared
4. MONITOR: Increase monitoring frequency to 1-minute intervals
5. ISOLATE: If malware confirmed, isolate from network

7.3 Capacity Planning Example

6-MONTH FORECAST:

Device Type: Server-LAB-03 (Database Server)

Current State (Nov 2024):

- Disk: 680 GB / 1 TB (68%)
- Memory: 48 GB / 64 GB (75%)
- CPU Cores: 16, avg utilization 35%

Projected State (May 2025):

- Disk: 950 GB / 1 TB (95%) ← CRITICAL
- Memory: 61 GB / 64 GB (95%) ← CRITICAL
- CPU: Adequate (35% avg utilization)

RECOMMENDATIONS:

Timeline	Action	Cost
----------	--------	------

Now	Add 500GB SSD	\$2,000
Month 3	Add 32GB RAM	\$800
Month 6	Upgrade to 2TB storage	\$3,500
Total 6-month cost: ~\$6,300		
Impact: Prevents 2-3 incidents, saves ~\$50k in downtime		

PHASE 8: Database Schema (Offline)

8.1 Core Tables

sql

-- Systems inventory

```
CREATE TABLE systems (
    system_id INTEGER PRIMARY KEY,
    hostname TEXT UNIQUE,
    ip_address TEXT,
    os_type TEXT,
    location TEXT,
    department TEXT,
    added_date TIMESTAMP,
    device_type TEXT, -- Server/Workstation/Laptop
    criticality TEXT -- Critical/High/Medium/Low
);
```

-- Time-series metrics (compressed)

```
CREATE TABLE metrics_timeseries (
    id INTEGER PRIMARY KEY,
    system_id INTEGER FOREIGN KEY,
    timestamp TIMESTAMP,
    cpu_avg FLOAT,
    cpu_max FLOAT,
    memory_used INTEGER,
    memory_total INTEGER,
    disk_used INTEGER,
    disk_total INTEGER,
    temp_cpu FLOAT,
    network_sent INTEGER,
    network_recv INTEGER,
    uptime_seconds INTEGER
);
```

-- Predictions

```
CREATE TABLE predictions (
    id INTEGER PRIMARY KEY,
    system_id INTEGER FOREIGN KEY,
    prediction_date TIMESTAMP,
    failure_probability FLOAT,
    risk_level TEXT, -- Critical/High/Medium/Low
    predicted_failure_window TEXT,
    primary_risk TEXT,
    model_version TEXT,
    confidence FLOAT
);
```

-- Anomalies

```
CREATE TABLE anomalies (
    id INTEGER PRIMARY KEY,
```

```
system_id INTEGER FOREIGN KEY,  
detection_timestamp TIMESTAMP,  
anomaly_type TEXT,  
anomaly_score FLOAT,  
affected_metrics TEXT, -- JSON  
root_cause_hypothesis TEXT,  
recommended_actions TEXT -- JSON  
);
```

-- Incidents

```
CREATE TABLE incidents (  
id INTEGER PRIMARY KEY,  
system_id INTEGER FOREIGN KEY,  
incident_date TIMESTAMP,  
severity TEXT,  
description TEXT,  
resolution TEXT,  
resolution_time_hours FLOAT,  
root_cause TEXT  
);
```

-- Alerts

```
CREATE TABLE alerts (  
id INTEGER PRIMARY KEY,  
system_id INTEGER FOREIGN KEY,  
alert_date TIMESTAMP,  
alert_type TEXT,  
severity TEXT,  
message TEXT,  
acknowledged BOOLEAN,  
acknowledged_by TEXT,  
acknowledged_date TIMESTAMP  
);
```

-- Model metadata

```
CREATE TABLE models (  
id INTEGER PRIMARY KEY,  
model_name TEXT,  
version TEXT,  
creation_date TIMESTAMP,  
training_samples INTEGER,  
accuracy FLOAT,  
precision FLOAT,  
recall FLOAT,  
f1_score FLOAT,
```

```
model_binary BLOB -- Serialized model
);
```

PHASE 9: Expected Outcomes for Enterprise Operations

9.1 Quantifiable Benefits

BEFORE (Manual Monitoring):

- Detection time: 24-48 hours (after failure)
- Downtime: 4-8 hours per incident
- Cost per incident: \$10,000-50,000
- Incidents per year: 8-12

AFTER (AI Monitoring System):

- Detection time: 24-48 hours BEFORE failure
- Downtime: 0.5-2 hours (planned maintenance)
- Cost per incident: \$1,000-5,000 (prevented incidents)
- Incidents per year: 1-2 (only unforeseeable)

ANNUAL ROI:

- Incidents prevented: $\sim 8 \times \$25,000 = \$200,000$ saved
- System cost: $\sim \$40,000$ (infrastructure)
- Net savings: $\$160,000+$ per year
- Payback period: ~ 3 months

9.2 Strategic Advantages

1. OPERATIONAL EXCELLENCE:

- ✓ 99.5%+ uptime (instead of 95%)
- ✓ Proactive maintenance (not reactive)
- ✓ Optimized resource allocation
- ✓ Compliance audit trail

2. SECURITY:

- ✓ Anomaly detection catches intrusions
- ✓ Unauthorized access alerts
- ✓ Network forensics support
- ✓ Air-gap safe (no external dependency)

3. PLANNING:

- ✓ Capacity planning with 6-month foresight
- ✓ Budget forecasting accuracy
- ✓ Lifecycle management optimization
- ✓ Vendor negotiation data

4. KNOWLEDGE:

- ✓ Understanding failure patterns
- ✓ Process improvement opportunities
- ✓ Risk assessment capability
- ✓ Training and documentation

PHASE 10: Quick Start Implementation

Step 1: Build Core Collector (Days 1-3)

Create `agent.py` on each device:

python

- Collect metrics every **5** minutes
- Store locally **in** SQLite
- Ready **for** master pull

Step 2: Build Master Aggregator (Days 4-7)

Create `master_collector.py`:

python

- SSH into each device
- Pull local SQLite data
- Consolidate into central DB
- Generate alerts

Step 3: Build ML Pipeline (Days 8-14)

Create `ml_pipeline.py`:

python

- Load historical data
- Feature engineering
- Train models (Random Forest, Isolation Forest)
- Save models **for** inference

Step 4: Build Dashboard (Days 15-21)

Create `dashboard.py`:

python

- Visualize all devices
- Show predictions
- Display anomalies
- Generate reports

Final Recommendations for Enterprise Project

What to Prioritize:

- 1. Phase 1 (Multi-Device):** This is your MVP. Get this working first.
- 2. Phase 2 (Data Storage):** Foundation for everything else.
- 3. Phase 4 (ML Basics):** Start with simple models (Random Forest, Isolation Forest) before complex ones.
- 4. Phase 5 (Security):** Critical for DRDO - encrypt everything.
- 5. Phase 7 (Reporting):** Stakeholders want clear reports.

What to Avoid:

- ✗ Don't try real-time ML on all devices (too heavy)
- ✗ Don't start with complex LSTM models
- ✗ Don't build web dashboard before core logic
- ✗ Don't mix online and offline mode
- ✓ Build it completely offline first

Success Metrics for Your Internship:

- Week 4: Multi-device monitoring functional (10+ devices)
- Week 8: ML predictions deployed (failure prediction model working)
- Week 12: Complete system with reports and automation

Present findings:

- How many devices monitored?
- How many potential failures predicted?
- How many incidents prevented?
- Cost savings quantified?
- Future roadmap proposed?

Conclusion

You're building a **critical infrastructure monitoring system** that can:

- Monitor 100s of devices in real-time

- Predict failures before they happen
- Detect security anomalies
- Generate actionable reports
- Run completely offline

This transforms DRDO from **reactive (fixing broken systems)** to **proactive (preventing failures)** — a massive operational upgrade.

Start simple, iterate fast, measure impact. Good luck! 